

d9vuyaffr

December 3, 2024

1 Car Price Prediction Model

1.1 Name: Ajo Babu A

1.2 1. Overview of Problem Statement:

A Chinese automobile company aspires to establish its presence in the US market by manufacturing cars locally to compete with US and European brands. Understanding the factors influencing car prices in the US, which differ from the Chinese market, is critical for success. The company has contracted a consulting firm to analyze a dataset of various cars from the American market, with the goal of identifying the variables that significantly affect car prices and understanding how these factors influence pricing. This insight will guide the company's market entry strategy and product design to achieve competitive pricing.

1.3 2. Objective

The objective of this project is to develop a regression model to predict car prices in the US market using various independent variables. This model will help identify significant factors influencing car prices and explain their relationship to pricing, enabling the company to design competitive cars and devise effective business strategies.

1.4 3. Data description

Source: https://drive.google.com/file/d/1FHmYNLs9v0Enc-UExEMpitOFGsWvB2dP/view?usp=drive_link

Features: car_ID ,symboling, CarName, fueltypes, aspiration, doornumber, carbody, drivewheel, enginelocation, wheelbase, carlength, carwidth, carheight, curbweight, enginetype, cylindernumber, enginesize, fuelsystem, boreratio, stroke, compressionratio, horsepower, peakrpm,citympg, highwaympg, price

1.5 Importing Libraries

```
[271]: # Importing Libraries

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.preprocessing import LabelEncoder, OneHotEncoder
```

```

from sklearn.feature_selection import VarianceThreshold

from sklearn.preprocessing import StandardScaler, MinMaxScaler

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
from sklearn.svm import SVR
from sklearn.metrics import r2_score, mean_absolute_error, mean_squared_error

from sklearn.model_selection import GridSearchCV

import joblib
import warnings
warnings.filterwarnings("ignore")

```

1.6 4. Importing Data

```

[20]: data = pd.read_csv('CarPrice_Assignment.csv')
data

```

```

[20]:
   car_ID  symboling      CarName fueltype aspiration \
0         1          3    alfa-romero giulia      gas      std
1         2          3    alfa-romero stelvio      gas      std
2         3          1  alfa-romero Quadrifoglio      gas      std
3         4          2          audi 100 ls      gas      std
4         5          2          audi 100ls      gas      std
..      ...      ...
200      201         -1    volvo 145e (sw)      gas      std
201      202         -1    volvo 144ea      gas    turbo
202      203         -1    volvo 244dl      gas      std
203      204         -1    volvo 246      diesel    turbo
204      205         -1    volvo 264gl      gas    turbo

   doornumber  carbody drivewheel enginelocation  wheelbase  ... \
0          two  convertible      rwd          front      88.6  ...
1          two  convertible      rwd          front      88.6  ...
2          two   hatchback      rwd          front      94.5  ...
3         four      sedan      fwd          front      99.8  ...
4         four      sedan      4wd          front      99.4  ...
..      ...      ...
200      four      sedan      rwd          front     109.1  ...
201      four      sedan      rwd          front     109.1  ...
202      four      sedan      rwd          front     109.1  ...
203      four      sedan      rwd          front     109.1  ...

```

204	four	sedan	rwd	front	109.1	...
-----	------	-------	-----	-------	-------	-----

	enginesize	fuelsystem	boreratio	stroke	compressionratio	horsepower	\
0	130	mpfi	3.47	2.68	9.0	111	
1	130	mpfi	3.47	2.68	9.0	111	
2	152	mpfi	2.68	3.47	9.0	154	
3	109	mpfi	3.19	3.40	10.0	102	
4	136	mpfi	3.19	3.40	8.0	115	
..	
200	141	mpfi	3.78	3.15	9.5	114	
201	141	mpfi	3.78	3.15	8.7	160	
202	173	mpfi	3.58	2.87	8.8	134	
203	145	idi	3.01	3.40	23.0	106	
204	141	mpfi	3.78	3.15	9.5	114	

	peakrpm	citympg	highwaympg	price
0	5000	21	27	13495.0
1	5000	21	27	16500.0
2	5000	19	26	16500.0
3	5500	24	30	13950.0
4	5500	18	22	17450.0
..
200	5400	23	28	16845.0
201	5300	19	25	19045.0
202	5500	18	23	21485.0
203	4800	26	27	22470.0
204	5400	19	25	22625.0

[205 rows x 26 columns]

```
[22]: df = pd.DataFrame(data)
```

```
[24]: df.head()
```

```
[24]:
```

	car_ID	symboling	CarName	fueltype	aspiration	doornumber	\
0	1	3	alfa-romero giulia	gas	std	two	
1	2	3	alfa-romero stelvio	gas	std	two	
2	3	1	alfa-romero Quadrifoglio	gas	std	two	
3	4	2	audi 100 ls	gas	std	four	
4	5	2	audi 100ls	gas	std	four	

	carbody	drivewheel	engine location	wheelbase	...	enginesize	\
0	convertible	rwd	front	88.6	...	130	
1	convertible	rwd	front	88.6	...	130	
2	hatchback	rwd	front	94.5	...	152	
3	sedan	fwd	front	99.8	...	109	
4	sedan	4wd	front	99.4	...	136	

	fuelsystem	boreratio	stroke	compressionratio	horsepower	peakrpm	citympg \
0	mpfi	3.47	2.68	9.0	111	5000	21
1	mpfi	3.47	2.68	9.0	111	5000	21
2	mpfi	2.68	3.47	9.0	154	5000	19
3	mpfi	3.19	3.40	10.0	102	5500	24
4	mpfi	3.19	3.40	8.0	115	5500	18

	highwaympg	price
0	27	13495.0
1	27	16500.0
2	26	16500.0
3	30	13950.0
4	22	17450.0

[5 rows x 26 columns]

```
[26]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
#   Column                Non-Null Count  Dtype
---  -
0   car_ID                205 non-null    int64
1   symboling              205 non-null    int64
2   CarName               205 non-null    object
3   fueltype              205 non-null    object
4   aspiration            205 non-null    object
5   doornumber            205 non-null    object
6   carbody              205 non-null    object
7   drivewheel           205 non-null    object
8   enginelocation        205 non-null    object
9   wheelbase            205 non-null    float64
10  carlength            205 non-null    float64
11  carwidth             205 non-null    float64
12  carheight            205 non-null    float64
13  curbweight           205 non-null    int64
14  enginetype           205 non-null    object
15  cylindernumber       205 non-null    object
16  enginesize           205 non-null    int64
17  fuelsystem           205 non-null    object
18  boreratio            205 non-null    float64
19  stroke               205 non-null    float64
20  compressionratio     205 non-null    float64
21  horsepower           205 non-null    int64
22  peakrpm              205 non-null    int64
```

```

23  citympg          205 non-null    int64
24  highwaympg      205 non-null    int64
25  price           205 non-null    float64
dtypes: float64(8), int64(8), object(10)
memory usage: 41.8+ KB

```

```
[28]: df.dtypes
```

```

[28]: car_ID          int64
      symboling      int64
      CarName        object
      fueltype       object
      aspiration     object
      doornumber     object
      carbody        object
      drivewheel     object
      enginelocation object
      wheelbase      float64
      carlength      float64
      carwidth       float64
      carheight      float64
      curbweight     int64
      enginetype     object
      cylindernumber object
      enginesize      int64
      fuelsystem     object
      boreratio      float64
      stroke         float64
      compressionratio float64
      horsepower     int64
      peakrpm        int64
      citympg        int64
      highwaympg     int64
      price          float64
      dtype: object

```

```
[30]: df.describe()
```

```

[30]:
count      car_ID      symboling      wheelbase      carlength      carwidth      carheight  \
count    205.000000    205.000000    205.000000    205.000000    205.000000    205.000000
mean      103.000000      0.834146      98.756585     174.049268      65.907805      53.724878
std        59.322565      1.245307       6.021776      12.337289       2.145204       2.443522
min         1.000000     -2.000000      86.600000     141.100000      60.300000      47.800000
25%        52.000000      0.000000      94.500000     166.300000      64.100000      52.000000
50%       103.000000      1.000000      97.000000     173.200000      65.500000      54.100000
75%       154.000000      2.000000     102.400000     183.100000      66.900000      55.500000
max       205.000000      3.000000     120.900000     208.100000      72.300000      59.800000

```

	curbweight	enginesize	boreratio	stroke	compressionratio \
count	205.000000	205.000000	205.000000	205.000000	205.000000
mean	2555.565854	126.907317	3.329756	3.255415	10.142537
std	520.680204	41.642693	0.270844	0.313597	3.972040
min	1488.000000	61.000000	2.540000	2.070000	7.000000
25%	2145.000000	97.000000	3.150000	3.110000	8.600000
50%	2414.000000	120.000000	3.310000	3.290000	9.000000
75%	2935.000000	141.000000	3.580000	3.410000	9.400000
max	4066.000000	326.000000	3.940000	4.170000	23.000000

	horsepower	peakrpm	citympg	highwaympg	price
count	205.000000	205.000000	205.000000	205.000000	205.000000
mean	104.117073	5125.121951	25.219512	30.751220	13276.710571
std	39.544167	476.985643	6.542142	6.886443	7988.852332
min	48.000000	4150.000000	13.000000	16.000000	5118.000000
25%	70.000000	4800.000000	19.000000	25.000000	7788.000000
50%	95.000000	5200.000000	24.000000	30.000000	10295.000000
75%	116.000000	5500.000000	30.000000	34.000000	16503.000000
max	288.000000	6600.000000	49.000000	54.000000	45400.000000

```
[34]: df.shape
```

```
[34]: (205, 26)
```

```
[36]: df.columns
```

```
[36]: Index(['car_ID', 'symboling', 'CarName', 'fueltype', 'aspiration',
        'doornumber', 'carbody', 'drivewheel', 'enginelocation', 'wheelbase',
        'carlength', 'carwidth', 'carheight', 'curbweight', 'enginetype',
        'cylindernumber', 'enginesize', 'fuelsystem', 'boreratio', 'stroke',
        'compressionratio', 'horsepower', 'peakrpm', 'citympg', 'highwaympg',
        'price'],
        dtype='object')
```

1.7 5.Data preprocessing and Data cleaning

```
[39]: # Checking for duplicate
df.duplicated()
```

```
[39]: 0      False
      1      False
      2      False
      3      False
      4      False
      ...
     200     False
```

```

201    False
202    False
203    False
204    False
Length: 205, dtype: bool

```

```
[43]: df.duplicated().sum()
```

```
[43]: 0
```

No duplicates found

```
[46]: # Checking for null values
df.isnull()
```

```
[46]:
```

	car_ID	symboling	CarName	fueltype	aspiration	doornumber	carbody	\
0	False	False	False	False	False	False	False	
1	False	False	False	False	False	False	False	
2	False	False	False	False	False	False	False	
3	False	False	False	False	False	False	False	
4	False	False	False	False	False	False	False	
..	
200	False	False	False	False	False	False	False	
201	False	False	False	False	False	False	False	
202	False	False	False	False	False	False	False	
203	False	False	False	False	False	False	False	
204	False	False	False	False	False	False	False	

	drivewheel	enginelocation	wheelbase	...	enginesize	fuelsystem	\
0	False	False	False	...	False	False	
1	False	False	False	...	False	False	
2	False	False	False	...	False	False	
3	False	False	False	...	False	False	
4	False	False	False	...	False	False	
..	
200	False	False	False	...	False	False	
201	False	False	False	...	False	False	
202	False	False	False	...	False	False	
203	False	False	False	...	False	False	
204	False	False	False	...	False	False	

	boreratio	stroke	compressionratio	horsepower	peakrpm	citympg	\
0	False	False	False	False	False	False	
1	False	False	False	False	False	False	
2	False	False	False	False	False	False	
3	False	False	False	False	False	False	
4	False	False	False	False	False	False	

..
200	False	False	False	False	False	False
201	False	False	False	False	False	False
202	False	False	False	False	False	False
203	False	False	False	False	False	False
204	False	False	False	False	False	False

	highwaympg	price
0	False	False
1	False	False
2	False	False
3	False	False
4	False	False

..
200	False	False
201	False	False
202	False	False
203	False	False
204	False	False

[205 rows x 26 columns]

```
[48]: df.isnull().sum()
```

```
[48]: car_ID          0
      symboling      0
      CarName        0
      fueltype       0
      aspiration     0
      doornumber     0
      carbody        0
      drivewheel     0
      enginelocation  0
      wheelbase      0
      carlength      0
      carwidth       0
      carheight      0
      curbweight     0
      enginetype     0
      cylindernumber  0
      enginesize      0
      fuelsystem     0
      boreratio      0
      stroke         0
      compressionratio 0
      horsepower     0
      peakrpm        0
```



```
citympg          0
highwaympg       0
price            0
dtype: int64
```

No null values found

1.7.1 Classifying columns


```
[52]: numerical_columns = df.select_dtypes(include=['number']).columns
      categorical_columns = df.select_dtypes(include=['object']).columns


      print("numerical columns: ", numerical_columns)
      print("Categorical columns: ", categorical_columns)
```


```
numerical columns: Index(['car_ID', 'symboling', 'wheelbase', 'carlength',
                          'carwidth',
                          'carheight', 'curbweight', 'enginesize', 'bore_ratio', 'stroke',
                          'compressionratio', 'horsepower', 'peakrpm', 'citympg', 'highwaympg',
                          'price'],
                          dtype='object')
Categorical columns: Index(['CarName', 'fueltype', 'aspiration', 'doornumber',
                          'carbody',
                          'drivewheel', 'engine_location', 'engine_type', 'cylindernumber',
                          'fuelsystem'],
                          dtype='object')
```


1.7.2 Outlier Detection


```
[57]: # boxplot to identify outliers
      for i in df.select_dtypes(include='number').columns:
          sns.boxplot(data=df, x=i)
          plt.show()
```













































```
[59]: numerical_columns = df.select_dtypes(include=['number']).columns
plt.figure(figsize=(10,6))
sns.boxplot(data = df[numerical_columns])
plt.xticks(rotation=90)
```

```
[59]: ([0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15],
      [Text(0, 0, 'car_ID'),
       Text(1, 0, 'symboling'),
       Text(2, 0, 'wheelbase'),
       Text(3, 0, 'carlength'),
       Text(4, 0, 'carwidth'),
       Text(5, 0, 'carheight'),
       Text(6, 0, 'curbweight'),
       Text(7, 0, 'enginesize'),
       Text(8, 0, 'boreratio'),
       Text(9, 0, 'stroke'),
       Text(10, 0, 'compressionratio'),
       Text(11, 0, 'horsepower'),
       Text(12, 0, 'peakrpm'),
       Text(13, 0, 'citympg'),
       Text(14, 0, 'highwaympg')],
```

```
Text(15, 0, 'price']])
```



Outliers Found

1.7.3 IQR Method

```
[64]: # List of features with potential outliers
features = ['wheelbase', 'carlength', 'carwidth', 'enginesize', 'stroke',
            'compressionratio', 'horsepower', 'peakrpm', 'citympg',
            ↪ 'highwaympg', 'price']

# Function to apply IQR method to fix outliers
def fix_outliers_iqr(df, columns):
    for col in columns:
        Q1 = df[col].quantile(0.25)
        Q3 = df[col].quantile(0.75)
        IQR = Q3 - Q1

        lower_bound = Q1 - 1.5 * IQR
        upper_bound = Q3 + 1.5 * IQR

    # Replace outliers with the respective bounds
```

```

        df[col] = df[col].apply(lambda x: lower_bound if x < lower_bound else
↪upper_bound if x > upper_bound else x)

    return df

```


```
df = fix_outliers_iqr(df, features)
```


1.7.4 Visualising after Outlier detection


```


[67]: # boxplot to identify outliers
for i in df.select_dtypes(include='number').columns:
    sns.boxplot(data=df, x=i)
    plt.show()


```













































```
[76]: numerical_columns = df.select_dtypes(include=['number']).columns
plt.figure(figsize=(15,8))
sns.boxplot(data = df[numerical_columns])
plt.xticks(rotation=90)
```


```
[76]: ([0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15],
      [Text(0, 0, 'car_ID'),
       Text(1, 0, 'symboling'),
       Text(2, 0, 'wheelbase'),
       Text(3, 0, 'carlength'),
       Text(4, 0, 'carwidth'),
       Text(5, 0, 'carheight'),
       Text(6, 0, 'curbweight'),
       Text(7, 0, 'enginesize'),
       Text(8, 0, 'boreratio'),
       Text(9, 0, 'stroke'),
       Text(10, 0, 'compressionratio'),
       Text(11, 0, 'horsepower'),
       Text(12, 0, 'peakrpm'),
       Text(13, 0, 'citympg'),
       Text(14, 0, 'highwaympg')],
```

```
Text(15, 0, 'price']])
```



1.8 Checking Skew

```
[90]: # visualising
plt.figure(figsize=(12,6))
sns.histplot(df[numerical_columns])
plt.show()
```




```
[84]: df[numerical_columns].skew()
```

```
[84]: car_ID          0.000000
      symboling      0.211072
      wheelbase     0.924916
      carlength     0.155954
      carwidth      0.776278
      carheight     0.063123
      curbweight    0.681398
      enginesize     0.908453
      boreratio     0.020156
      stroke        -0.379130
      compressionratio 0.035149
      horsepower    0.814957
      peakrpm       0.049935
      citympg       0.604594
      highwaympg    0.347441
      price         1.222031
      dtype: float64
```

```
[92]: sns.distplot(df['price'])
```

```
[92]: <Axes: xlabel='price', ylabel='Density'>
```



Skewness in high for 'price'

1.8.1 Fixing skewness using log transformation

```
[99]: df['price'] = np.log(df['price'])
df['wheelbase'] = np.log(df['wheelbase'])
df['carwidth'] = np.log(df['carwidth'])
df['enginesize'] = np.log(df['enginesize'])
df['compressionratio'] = np.log(df['compressionratio'])
df['horsepower'] = np.log(df['horsepower'])
```

```
[107]: df[numerical_columns].skew()
```

```
[107]: car_ID          0.000000
symboling          0.211072
wheelbase          0.781960
carlength          0.155954
carwidth           0.696287
carheight          0.063123
curbweight         0.681398
```

```

enginesize      0.401418
boreratio       0.020156
stroke          -0.379130
compressionratio -0.241483
horsepower      0.287093
peakrpm         0.049935
citympg         0.604594
highwaympg      0.347441
price           0.459254
dtype: float64

```

```
[109]: df['compressionratio'] = np.sqrt(df['compressionratio'])
df['wheelbase'] = np.sqrt(df['wheelbase'])
```


```
[111]: df[numerical_columns].skew()
```

```
[111]: car_ID      0.000000
symboling      0.211072
wheelbase      0.766142
carlength      0.155954
carwidth       0.696287
carheight      0.063123
curbweight     0.681398
enginesize     0.401418
boreratio      0.020156
stroke         -0.379130
compressionratio -0.304199
horsepower     0.287093
peakrpm        0.049935
citympg        0.604594
highwaympg     0.347441
price          0.459254
dtype: float64

```

```
[113]: sns.distplot(df['price'])
```

```
[113]: <Axes: xlabel='price', ylabel='Density'>
```



```
[115]: df[numerical_columns].skew()
```

```
[115]: car_ID          0.000000
       symboling      0.211072
       wheelbase     0.766142
       carlength     0.155954
       carwidth       0.696287
       carheight      0.063123
       curbweight     0.681398
       enginesize     0.401418
       boreratio      0.020156
       stroke        -0.379130
       compressionratio -0.304199
       horsepower     0.287093
       peakrpm        0.049935
       citympg        0.604594
       highwaympg     0.347441
       price          0.459254
       dtype: float64
```

1.9 6. EDA

```
[118]: df1 = df.copy()
```

```
[120]: df1.head()
```

```
[120]:
```

	car_ID	symboling	CarName	fueltype	aspiration	doornumber	\
0	1	3	alfa-romero giulia	gas	std	two	
1	2	3	alfa-romero stelvio	gas	std	two	
2	3	1	alfa-romero Quadrifoglio	gas	std	two	
3	4	2	audi 100 ls	gas	std	four	
4	5	2	audi 100ls	gas	std	four	

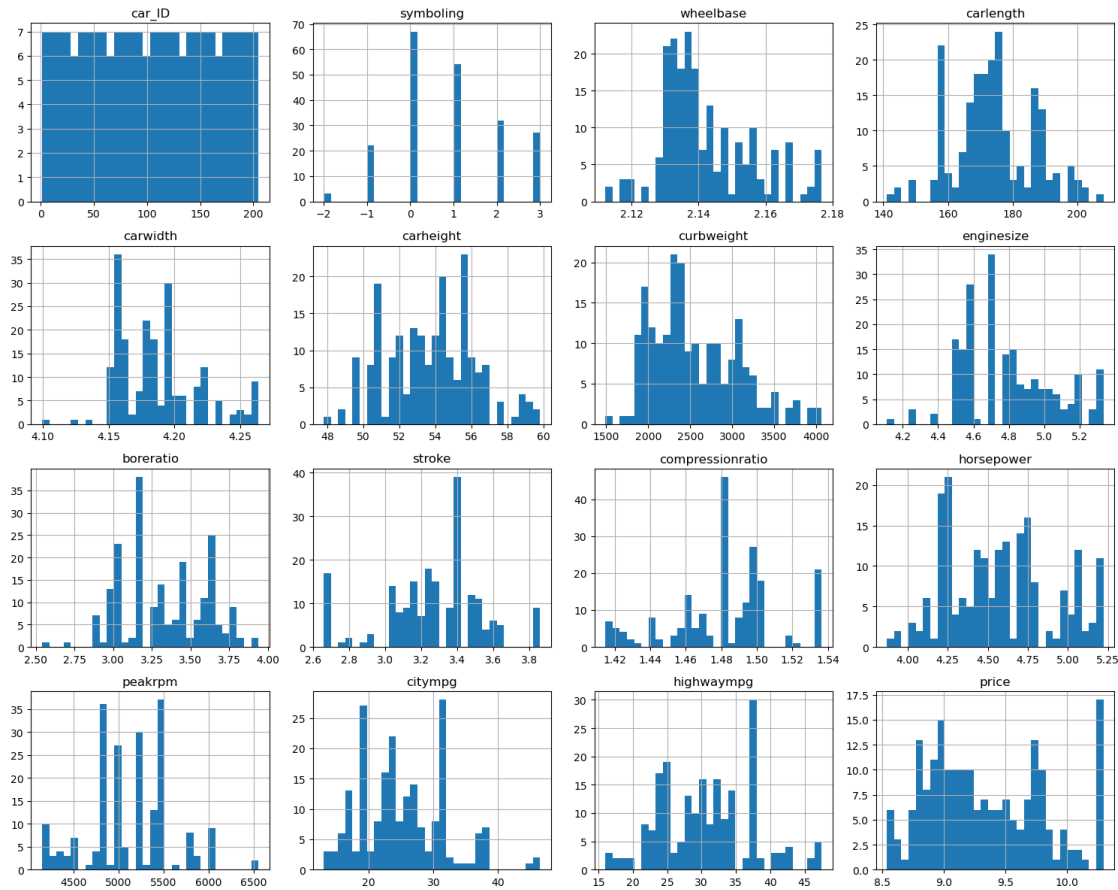
	carbody	drivewheel	engine	location	wheelbase	...	engine	size	\
0	convertible	rwd		front	2.117577	...		4.867534	
1	convertible	rwd		front	2.117577	...		4.867534	
2	hatchback	rwd		front	2.132745	...		5.023881	
3	sedan	fwd		front	2.145500	...		4.691348	
4	sedan	4wd		front	2.144563	...		4.912655	

	fuelsystem	boreratio	stroke	compressionratio	horsepower	peakrpm	citympg	\
0	mpfi	3.47	2.68	1.482304	4.709530	5000.0	21.0	
1	mpfi	3.47	2.68	1.482304	4.709530	5000.0	21.0	
2	mpfi	2.68	3.47	1.482304	5.036953	5000.0	19.0	
3	mpfi	3.19	3.40	1.517427	4.624973	5500.0	24.0	
4	mpfi	3.19	3.40	1.442027	4.744932	5500.0	18.0	

	highwaympg	price
0	27.0	9.510075
1	27.0	9.711116
2	26.0	9.711116
3	30.0	9.543235
4	22.0	9.767095


[5 rows x 26 columns]

```
[122]: # Histogram
df1.hist(bins=30, figsize=(15, 12))
plt.tight_layout()
plt.show()
```


```
[132]: #Compute correlation matrix
corr_matrix = df1[numerical_columns].corr()



# Heatmap
plt.figure(figsize=(20, 22))
sns.heatmap(corr_matrix, annot=True, fmt=".2f", cmap="coolwarm")
plt.title("Correlation Heatmap")
plt.show()
```







```
[172]: # Plot distribution for numerical column
for feature in numerical_columns:
    plt.figure(figsize=(6, 4))
    sns.histplot(df1[feature], kde=True, bins=30, color='blue')
    plt.title(f'Distribution of {feature}')
```



```
plt.xlabel(feature)
plt.ylabel('Frequency')
plt.show()
```


































```
[158]: categorical_feature = df1.drop(numerical_columns, axis=1)

# Plot bar chart for each categorical feature
for feature in categorical_feature:
    plt.figure(figsize=(12, 6))
    sns.countplot(x=feature, data=df, palette='viridis')
    plt.title(f'Distribution of {feature}')
    plt.xlabel(feature)
    plt.ylabel('Count')
    plt.xticks(rotation=90)
    plt.show()
```










```
[174]: # Scatter plots for price vs numerical features
features_to_compare = ['engine_size', 'horsepower', 'curbweight', 'citympg',
                        ↪ 'highwaympg']



for feature in features_to_compare:
```



```
plt.figure(figsize=(6, 4))
sns.scatterplot(x=df[feature], y=df['price'], color='green')
plt.title(f'Price vs {feature}')
plt.xlabel(feature)
plt.ylabel('Price')
plt.show()
```







1.10 7. Feature Engineering

Label encoding and Onehot Encoding

```
[183]: label_encoder = LabelEncoder()
df_label = df1.copy()
df1.head()
```

```
[183]:   car_ID  symboling      CarName fueltype aspiration doornumber \
0        1          3  alfa-romero giulia      gas         std         two
1        2          3  alfa-romero stelvio      gas         std         two
2        3          1  alfa-romero Quadrifoglio      gas         std         two
3        4          2          audi 100 ls      gas         std         four
4        5          2          audi 100ls      gas         std         four

      carbody drivewheel enginelocation  wheelbase  ...  enginesize  \
0  convertible      rwd          front   2.117577  ...    4.867534
1  convertible      rwd          front   2.117577  ...    4.867534
2   hatchback      rwd          front   2.132745  ...    5.023881
3        sedan      fwd          front   2.145500  ...    4.691348
4        sedan      4wd          front   2.144563  ...    4.912655

      fuelsystem  boreratio  stroke  compressionratio  horsepower  peakrpm  citympg  \
0         mpfi        3.47    2.68          1.482304    4.709530    5000.0    21.0
1         mpfi        3.47    2.68          1.482304    4.709530    5000.0    21.0
2         mpfi        2.68    3.47          1.482304    5.036953    5000.0    19.0
3         mpfi        3.19    3.40          1.517427    4.624973    5500.0    24.0
4         mpfi        3.19    3.40          1.442027    4.744932    5500.0    18.0

      highwaympg      price
0          27.0  9.510075
1          27.0  9.711116
2          26.0  9.711116
3          30.0  9.543235
4          22.0  9.767095
```

[5 rows x 26 columns]

```
[185]: #label Encod
df_label['fueltype'] = label_encoder.fit_transform(df1['fueltype'])
df_label['aspiration'] = label_encoder.fit_transform(df1['aspiration'])
df_label['doornumber'] = label_encoder.fit_transform(df1['doornumber'])
df_label['carbody'] = label_encoder.fit_transform(df1['carbody'])
df_label['drivewheel'] = label_encoder.fit_transform(df1['drivewheel'])
df_label['enginelocation'] = label_encoder.fit_transform(df1['enginelocation'])
df_label['enginetype'] = label_encoder.fit_transform(df1['enginetype'])
df_label['cylindernumber'] = label_encoder.fit_transform(df1['cylindernumber'])
df_label['fuelsystem'] = label_encoder.fit_transform(df1['fuelsystem'])
```

```
df_label.head()
```

```
[185]:
```

	car_ID	symboling	CarName	fueltype	aspiration	\
0	1	3	alfa-romero giulia	1	0	
1	2	3	alfa-romero stelvio	1	0	
2	3	1	alfa-romero Quadrifoglio	1	0	
3	4	2	audi 100 ls	1	0	
4	5	2	audi 100ls	1	0	

	doornumber	carbody	drivewheel	enginelocation	wheelbase	...	\
0	1	0	2	0	2.117577	...	
1	1	0	2	0	2.117577	...	
2	1	2	2	0	2.132745	...	
3	0	3	1	0	2.145500	...	
4	0	3	0	0	2.144563	...	

	enginesize	fuelsystem	boreratio	stroke	compressionratio	horsepower	\
0	4.867534	5	3.47	2.68	1.482304	4.709530	
1	4.867534	5	3.47	2.68	1.482304	4.709530	
2	5.023881	5	2.68	3.47	1.482304	5.036953	
3	4.691348	5	3.19	3.40	1.517427	4.624973	
4	4.912655	5	3.19	3.40	1.442027	4.744932	

	peakrpm	citympg	highwaympg	price
0	5000.0	21.0	27.0	9.510075
1	5000.0	21.0	27.0	9.711116
2	5000.0	19.0	26.0	9.711116
3	5500.0	24.0	30.0	9.543235
4	5500.0	18.0	22.0	9.767095

[5 rows x 26 columns]

```
[187]: # OneHot encoding

onehot = OneHotEncoder(sparse_output=False)
hot_encod = onehot.fit_transform(df_label[['CarName']])
hot_columns = onehot.get_feature_names_out(['CarName'])
```

```
[189]: # Creating DataFrame with onehot encoded columns

df_onehot = pd.concat([
    df_label,
    pd.DataFrame(hot_encod, columns=hot_columns)
], axis=1)
```

```
[191]: df_onehot.head()
```

```

[191]:   car_ID  symboling          CarName  fueltype  aspiration  \
0      1      3      alfa-romero giulia      1      0
1      2      3      alfa-romero stelvio      1      0
2      3      1  alfa-romero Quadrifoglio      1      0
3      4      2      audi 100 ls      1      0
4      5      2      audi 100ls      1      0

      doornumber  carbody  drivewheel  enginelocation  wheelbase  ...  \
0      1      0      2      0  2.117577  ...
1      1      0      2      0  2.117577  ...
2      1      2      2      0  2.132745  ...
3      0      3      1      0  2.145500  ...
4      0      3      0      0  2.144563  ...

      CarName_volkswagen type 3  CarName_volvo 144ea  CarName_volvo 145e (sw)  \
0      0.0      0.0      0.0
1      0.0      0.0      0.0
2      0.0      0.0      0.0
3      0.0      0.0      0.0
4      0.0      0.0      0.0

      CarName_volvo 244dl  CarName_volvo 245  CarName_volvo 246  \
0      0.0      0.0      0.0
1      0.0      0.0      0.0
2      0.0      0.0      0.0
3      0.0      0.0      0.0
4      0.0      0.0      0.0

      CarName_volvo 264gl  CarName_volvo diesel  CarName_vw dasher  \
0      0.0      0.0      0.0
1      0.0      0.0      0.0
2      0.0      0.0      0.0
3      0.0      0.0      0.0
4      0.0      0.0      0.0

      CarName_vw rabbit
0      0.0
1      0.0
2      0.0
3      0.0
4      0.0

```

[5 rows x 173 columns]

```
[195]: df_onehot.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Columns: 173 entries, car_ID to CarName_vw rabbit
dtypes: float64(160), int32(9), int64(3), object(1)
memory usage: 270.0+ KB
```

```
[197]: # Dropping categorical values
df_onehot = df_onehot.drop('CarName',axis=1)
```

```
[199]: df_onehot.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Columns: 172 entries, car_ID to CarName_vw rabbit
dtypes: float64(160), int32(9), int64(3)
memory usage: 268.4 KB
```

1.10.1 Setting y as target

```
[204]: y = df_onehot['price']
y
x = df_onehot.drop(['price','car_ID'], axis=1)
x
```

```
[204]:
```

	symboling	fueltype	aspiration	doornumber	carbody	drivewheel	\
0	3	1	0	1	0	2	
1	3	1	0	1	0	2	
2	1	1	0	1	2	2	
3	2	1	0	0	3	1	
4	2	1	0	0	3	0	
..	
200	-1	1	0	0	3	2	
201	-1	1	1	0	3	2	
202	-1	1	0	0	3	2	
203	-1	0	1	0	3	2	
204	-1	1	1	0	3	2	

	enginelocation	wheelbase	carlength	carwidth	...	\
0	0	2.117577	168.8	4.160444	...	
1	0	2.117577	168.8	4.160444	...	
2	0	2.132745	171.2	4.182050	...	
3	0	2.145500	176.6	4.192680	...	
4	0	2.144563	176.6	4.195697	...	
..	
200	0	2.166164	188.8	4.232656	...	
201	0	2.166164	188.8	4.231204	...	
202	0	2.166164	188.8	4.232656	...	

203	0	2.166164	188.8	4.232656	...
204	0	2.166164	188.8	4.232656	...

	CarName_volkswagen type 3	CarName_volvo 144ea	CarName_volvo 145e (sw)	\
0	0.0	0.0	0.0	
1	0.0	0.0	0.0	
2	0.0	0.0	0.0	
3	0.0	0.0	0.0	
4	0.0	0.0	0.0	
..	
200	0.0	0.0	1.0	
201	0.0	1.0	0.0	
202	0.0	0.0	0.0	
203	0.0	0.0	0.0	
204	0.0	0.0	0.0	

	CarName_volvo 244dl	CarName_volvo 245	CarName_volvo 246	\
0	0.0	0.0	0.0	
1	0.0	0.0	0.0	
2	0.0	0.0	0.0	
3	0.0	0.0	0.0	
4	0.0	0.0	0.0	
..	
200	0.0	0.0	0.0	
201	0.0	0.0	0.0	
202	1.0	0.0	0.0	
203	0.0	0.0	1.0	
204	0.0	0.0	0.0	

	CarName_volvo 264gl	CarName_volvo diesel	CarName_vw dasher	\
0	0.0	0.0	0.0	
1	0.0	0.0	0.0	
2	0.0	0.0	0.0	
3	0.0	0.0	0.0	
4	0.0	0.0	0.0	
..	
200	0.0	0.0	0.0	
201	0.0	0.0	0.0	
202	0.0	0.0	0.0	
203	0.0	0.0	0.0	
204	1.0	0.0	0.0	

	CarName_vw rabbit
0	0.0
1	0.0
2	0.0
3	0.0


```

4          0.0
..        ...
200        0.0
201        0.0
202        0.0
203        0.0
204        0.0

```

[205 rows x 170 columns]

1.11 8. Feature Selection

1.11.1 Variance threshold

```

[208]: var_threshold = VarianceThreshold(threshold=0.5)
x_var = var_threshold.fit_transform(x)
var_selected = x.columns[var_threshold.get_support()].tolist()
print("1.Filter method results: ")
print("\na) Variance threshold")
print(f"Features selected: {len(var_selected)}")
print("Selected Features: ", var_selected[:12], "....")

```

1.Filter method results:

a) Variance threshold

Features selected: 11

Selected Features: ['symboling', 'carbody', 'carlength', 'carheight', 'curbweight', 'enginetype', 'cylindernumber', 'fuelsystem', 'peakrpm', 'citympg', 'highwaympg'] ...

```

[212]: var_selected

```

```

[212]: ['symboling',
'carbody',
'carlength',
'carheight',
'curbweight',
'enginetype',
'cylindernumber',
'fuelsystem',
'peakrpm',
'citympg',
'highwaympg']

```

1.12 9. Feature Scaling

```
[220]: standard_scaler = StandardScaler()  
       minmax_scaler = MinMaxScaler()
```


```
[222]: # Applying scaling method  
       x_standardized = standard_scaler.fit_transform(x_var)
```

```
[226]: # Converting to DataFrame  
       df_std = pd.DataFrame(x_standardized)  
       df_std.head()
```

```
[226]:
```

	0	1	2	3	4	5	6	\
0	1.743470	-3.050975	-0.426521	-2.020417	-0.014566	-2.865105	-0.147475	
1	1.743470	-3.050975	-0.426521	-2.020417	-0.014566	-2.865105	-0.147475	
2	0.133509	-0.717207	-0.231513	-0.543527	0.514882	1.886890	1.112210	
3	0.938490	0.449677	0.207256	0.235942	-0.420797	-0.013908	-0.147475	
4	0.938490	0.449677	0.207256	0.235942	0.516807	-0.013908	-1.407161	
	7	8	9	10				
0	0.869568	-0.262757	-0.649321	-0.552143				
1	0.869568	-0.262757	-0.649321	-0.552143				
2	0.869568	-0.262757	-0.958163	-0.702161				
3	0.869568	0.791357	-0.186058	-0.102086				
4	0.869568	0.791357	-1.112584	-1.302237				

```
[230]: # Drawing Correlation  
       correlation = df_std.corr()  
       plt.figure(figsize=(15, 12))  
       sns.heatmap(correlation, annot=True, cmap='coolwarm')  
       plt.show()
```



1.13 10. Splitting Data to training and testing

```
[235]: X_train, X_test, y_train, y_test = train_test_split(df_std, y, test_size=0.2,
↳ random_state=42)
```

1.13.1 Linear Regression

Linear Regression is a simple and interpretable model that establishes a linear relationship between the dependent variable (target) and one or more independent variables (features). It minimizes the sum of squared residuals to find the best-fitting line. While effective for datasets with linear trends, it may underperform when relationships are complex or non-linear.

1.13.2 Decision Tree Regressor

The Decision Tree Regressor is a tree-based model that splits the dataset into subsets based on feature values, creating a tree structure. Each leaf represents a prediction, while branches signify decisions. It's intuitive and can capture non-linear relationships but may overfit the data, requiring pruning or other regularization techniques.

1.13.3 Random Forest Regressor

Random Forest Regressor is an ensemble learning method that combines multiple decision trees (trained on random subsets of the data and features) to improve prediction accuracy. It reduces overfitting and increases robustness but may be computationally expensive for large datasets.

1.13.4 Gradient Boosting Regressor

Gradient Boosting Regressor is another ensemble technique that builds sequential decision trees, where each tree corrects the errors of the previous one. It is highly effective for complex, non-linear relationships and offers flexibility through hyperparameter tuning. However, it is computationally intensive and sensitive to overfitting if not properly tuned.

1.13.5 Support Vector Regressor (SVR)

SVR uses the principles of Support Vector Machines to find a hyperplane that best fits the data within a margin of tolerance. It is powerful for small- to medium-sized datasets with non-linear relationships, especially when paired with kernels. However, it may struggle with large datasets and requires careful parameter selection for optimal performance.

1.14 11. Building Models

```
[238]: # Define models
models = {
    "Linear Regression": LinearRegression(),
    "Decision Tree Regressor": DecisionTreeRegressor(random_state=42),
    "Random Forest Regressor": RandomForestRegressor(random_state=42),
    "Gradient Boosting Regressor": GradientBoostingRegressor(random_state=42),
    "Support Vector Regressor": SVR(kernel='rbf')
}
```

```
[240]: # Train and evaluate models
results = {}

for name, model in models.items():
    # Train the model
    model.fit(X_train, y_train)

    # Predict on the test set
    y_pred = model.predict(X_test)

    # Calculate metrics
    r2 = r2_score(y_test, y_pred)
    mae = mean_absolute_error(y_test, y_pred)
    rmse = np.sqrt(mean_squared_error(y_test, y_pred))

    # Store results
    results[name] = {
        "R2 Score": r2,
```

```

    "MAE": mae,
    "RMSE": rmse
}

```

1.15 12. Model Evaluation Result

```

[243]: # Display results
results_df = pd.DataFrame(results).T
print(results_df)

```

	R ² Score	MAE	RMSE
Linear Regression	0.825697	0.161574	0.201772
Decision Tree Regressor	0.744972	0.170912	0.244063
Random Forest Regressor	0.876245	0.136639	0.170016
Gradient Boosting Regressor	0.891191	0.133739	0.159419
Support Vector Regressor	0.790854	0.177862	0.221020

1.16 Visualisation

```

[253]: # Function to plot actual vs predicted prices
def plot_actual_vs_predicted(y_test, y_pred, model_name):
    plt.figure(figsize=(10, 6))
    plt.scatter(y_test, y_pred, color='blue', alpha=0.6, label='Predicted vs Actual')
    plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], color='red', linestyle='--', label='Perfect Prediction')
    plt.title(f'Actual vs Predicted Prices: {model_name}')
    plt.xlabel('Actual Price')
    plt.ylabel('Predicted Price')
    plt.legend()
    plt.grid(True)
    plt.show()


```



```


[255]: # Iterate through models to plot predictions
for name, model in models.items():
    # Generate predictions for the test set
    y_pred = model.predict(X_test)

    # Call the plotting function
    plot_actual_vs_predicted(y_test, y_pred, name)

```







1.17 Finding the best model

```
[257]: # Find the best model based on R2 Score
best_model_name = results_df['R2 Score'].idxmax()
best_model_metrics = results_df.loc[best_model_name]

print(f"Best Model: {best_model_name}")
print("\nPerformance Metrics:")
print(best_model_metrics)
```

Best Model: Gradient Boosting Regressor

Performance Metrics:

R² Score 0.891191

MAE 0.133739

RMSE 0.159419

Name: Gradient Boosting Regressor, dtype: float64

```
[259]: # Rank models by R2 Score
ranked_models = results_df.sort_values(by='R2 Score', ascending=False)
print("Ranked Models by R2 Score:")
print(ranked_models)
```

Ranked Models by R² Score:

R ² Score	MAE	RMSE
----------------------	-----	------

Gradient Boosting Regressor	0.891191	0.133739	0.159419
Random Forest Regressor	0.876245	0.136639	0.170016
Linear Regression	0.825697	0.161574	0.201772
Support Vector Regressor	0.790854	0.177862	0.221020
Decision Tree Regressor	0.744972	0.170912	0.244063

1.17.1 Best Model is Gradient Boosting Regressor

1.18 13. Hyperparameter Tuning

```
[266]: # Define the model
gbr = GradientBoostingRegressor(random_state=42)

# Define hyperparameters to tune
param_grid = {
    'n_estimators': [100, 200, 300],
    'learning_rate': [0.01, 0.05, 0.1, 0.2],
    'max_depth': [3, 4, 5, 6],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4],
    'subsample': [0.6, 0.8, 1.0]
}

# Perform Grid Search with Cross Validation
grid_search = GridSearchCV(estimator=gbr, param_grid=param_grid,
                           scoring='r2', cv=5, verbose=2, n_jobs=-1)

# Fit the model
grid_search.fit(X_train, y_train)

# Best parameters and performance
print("Best Parameters:", grid_search.best_params_)
print("Best R² Score:", grid_search.best_score_)
```

Fitting 5 folds for each of 1296 candidates, totalling 6480 fits
 Best Parameters: {'learning_rate': 0.05, 'max_depth': 6, 'min_samples_leaf': 2, 'min_samples_split': 2, 'n_estimators': 200, 'subsample': 0.6}
 Best R² Score: 0.8970872363382394

```
[268]: # Use the best estimator for predictions
best_gbr = grid_search.best_estimator_
y_pred = best_gbr.predict(X_test)

# Evaluate on the test set
from sklearn.metrics import r2_score, mean_absolute_error, mean_squared_error
r2 = r2_score(y_test, y_pred)
mae = mean_absolute_error(y_test, y_pred)
rmse = np.sqrt(mean_squared_error(y_test, y_pred))
```

```
print("\nTest Set Performance:")
print(f"R2 Score: {r2:.4f}")
print(f"MAE: {mae:.2f}")
print(f"RMSE: {rmse:.2f}")
```

Test Set Performance:
R² Score: 0.9089
MAE: 0.12
RMSE: 0.15

1.19 Saving the Model

```
[273]: # Save the model to a file
joblib.dump(best_gbr, 'car_price_prediction_gb_model.joblib')
print("Model saved as 'car_price_prediction_gb_model.joblib'")
```

Model saved as 'car_price_prediction_gb_model.joblib'

1.20 Conclusion

Final Insights: The Gradient Boosting Regressor model provides actionable insights into the car price dynamics, which can help the business:

Design cars that meet specific price targets. Strategize pricing based on critical features influencing the price. This project demonstrates a structured approach to solving regression problems and highlights the value of machine learning in deriving business insights.

[]: